

Pointers and Arrays

Dr Deepayan Bhowmik Dr Andrea Bracciali

In this lecture

- ▶ Types: data + operations
 - ▶ Pointers and arrays
 - ▶ Pointer arithmetic
 - ▶ Arrays as arguments
-
- ▶ Multi-dimensional arrays (pointers)

Extras:

- ▶ Variable-length arrays (pointers)
-



Types

Name

int

Set of data

... -2, -1, 0, 1, ... max_int

Operations

$10 / 3 = 3$ (vs. $10.0 / 3.0 = 3.333\dots$)

Memory representation

00000000

00000000

00000000

00000000



Introduction

- ▶ In C, pointers, as a type, have their own **operations**.
- ▶ In particular, we will look at arithmetic on pointers – addition and subtraction – for **pointers to array elements**.
- ▶ This leads to an alternative way of processing arrays, in which **pointers take the place of array indexes**.
- ▶ The relationship between pointers and arrays in C is a close one, understanding which is critical for mastering C.
- ▶ This also gives you an idea of a “*memory-aware*” way of programming




Pointer Arithmetic

- ▶ We know that pointers can point to array elements:

```
int a[10], *p;  
p = &a[0];
```

- ▶ A graphical representation:




Pointer Arithmetic

- ▶ We can now access `a[0]` through `p`; for example, we can store the value 5 in `a[0]` by writing

`*p = 5;`

- ▶ An updated picture



Pointer Arithmetic

- ▶ If p points to an element of an array a , what could be the meaning of $p+1$?

Pointer Arithmetic

- ▶ If p points to an element of an array a , what could be the meaning of $p+1$?
- ▶ Other elements of a can be accessed by performing *pointer arithmetic* (or *address arithmetic*) on p .
- ▶ C supports three (and only three) forms of pointer arithmetic:
 1. Adding an integer to a pointer
 2. Subtracting an integer from a pointer
 3. Subtracting one pointer from another

1. Adding an Integer to a Pointer

- ▶ Assume that the following declarations are in effect:

```
int a[10], *p, *q, i;
```

- ▶ Adding an integer j to a pointer p yields
a pointer to the element j places after
the one that p points to.

- ▶ More precisely, if p points to the array element
 $a[i]$,
then $p + j$ points to
 $a[i+j]$.




1. Adding an Integer to a Pointer

- ▶ Example of pointer addition:

```
p = &a[2];
```

```
q = p + 3;
```

```
p += 6;
```




2. Subtracting an Integer from a Pointer


► If p points to $a[i]$, then $p - j$ points to $a[i - j]$.

► Example:


```
p = &a[8];
```



```
q = p - 3;
```



```
p -= 6;
```



3. Subtracting One Pointer from Another


- ▶ When one pointer is subtracted from another, **the result is the distance** (measured in **array elements**) between the pointers.
- ▶ If p points to $a[i]$ and q points to $a[j]$, then $p - q$ is equal to $i - j$.
- ▶ Example:

```
p = &a[5];
```

```
q = &a[1];
```

```
i = p - q;    /* i is 4
```

```
i = q - p;    /* i is -4
```



3. Subtracting One Pointer from Another

- ▶ When one pointer is subtracted from another, **the result is the distance** (measured in **array elements**) between the pointers.

- ▶ If p points to


$a[i],$

and q points to

$a[j],$

then $p - q$ is equal to

$i - j.$



Subtracting One Pointer from Another

Careful !

- ▶ Operations that cause undefined behavior:
 - ▶ Performing arithmetic on a pointer that
doesn't point to an array element
 - ▶ Subtracting pointers unless
both point to elements of the same array

Comparing Pointers

- ▶ Pointers can be compared using the relational operators ($<$, $<=$, $>$, $>=$) and the equality operators ($==$ and $!=$).
 - ▶ Using relational operators is **meaningful only for pointers to elements of the same array.**
- ▶ The outcome of the comparison depends on the relative positions of the two elements in the array.
- ▶ After the assignments

```
p = &a[5];  
q = &a[1];
```

the value of $p <= q$ is 0 (false) and
the value of $p >= q$ is 1 (true).

Pointers to Compound Literals (C99)

- ▶ It's legal for a pointer to point to an element within an array created by a compound literal (i.e., an object that has no name):

```
int *p = (int []){3, 0, 3, 4, 1};
```

- ▶ Using a compound literal saves us the trouble of first declaring an array variable and then making `p` point to the first element of that array:

```
int a[] = {3, 0, 3, 4, 1};  
int *p = &a[0];
```


Using Pointers for Array Processing

- ▶ Pointer arithmetic allows us to visit the elements of an array by repeatedly incrementing a pointer variable.
- ▶ A loop that sums the elements of an array `a`:

```
#define N 10
```

```
...
```

```
int a[N], sum, *p;
```

```
...
```


```
sum = 0;
```

```
for (p = &a[0]; p < &a[N]; p++)
```


```
    sum += *p;
```

Using Pointers for Array Processing


At the end of the first iteration:



At the end of the second iteration:



At the end of the third iteration:



Using Pointers for Array Processing

- ▶ The condition `p < &a[N]` in the `for` statement deserves special mention.
- ▶ It's legal to apply the address operator to `a[N]`, even though this element doesn't exist.
- ▶ Pointer arithmetic may save execution time.

Using an Array Name as a Pointer

- ▶ The name of an array can be used as a pointer to the first element in the array. E.g. with `int a[10];`

```
*a = 7;           /* stores 7 in a[0] */
*(a+1) = 12;      /* stores 12 in a[1] */
```

- ▶ In general, `a + i` is the same as `&a[i]`. Both represent a pointer to element `i` of `a`.
- ▶ Also, `*(a+i)` is equivalent to `a[i]`. Both represent element `i` itself.

Using an Array Name as a Pointer

- ▶ The fact that an array name can serve as a pointer makes it easier to write loops that step through an array.

- ▶ Original loop:

```
for (p = &a[0]; p < &a[N]; p++)  
    sum += *p;
```

- ▶ Simplified version:

```
for (p = a; p < a + N; p++)  
    sum += *p;
```

Using an Array Name as a Pointer

- ▶ Although an array name can be used as a pointer, it's **not possible to assign it a new value**.
- ▶ Attempting to make it point elsewhere is an error:

```
while (*a != 0)
    a++;                /* *** WRONG *** */
```

- ▶ This is no great loss; we can always copy `a` into a pointer variable, then change the pointer variable:

```
p = a;
while (*p != 0)
    p++;
```

Program: Reversing a Series of Numbers

- ▶ The `reverse.c` program reads 10 numbers, then writes the numbers in reverse order.
- ▶ A ‘naïve’ approach is to store the numbers in an array, with subscripting used to access elements of the array.
- ▶ `reverse3.c` is a new version of the program in which subscripting has been replaced with pointer arithmetic.

reverse3.c

```
/* Reverses a series of numbers (pointer version) */
#include <stdio.h>

#define N 10

int main(void)
{
    int a[N], *p;

    printf("Enter %d numbers: ", N);
    for (p = a; p < a + N; p++)
        scanf("%d", p); /* Recall: p is a ptr, so no & req'd */

    printf("In reverse order:");
    for (p = a + N - 1; p >= a; p--)
        printf(" %d", *p);
    printf("\n");

    return 0;
}
```



In this lecture

- ▶ Types: data + operations
 - ▶ Pointers and arrays
 - ▶ Pointer arithmetic
 - ▶ Arrays as arguments
-
- ▶ Multi-dimensional arrays (pointers)

Extras:

- ▶ Variable-length arrays (pointers)
-



Array Arguments (Revisited)

- ▶ When passed to a function, an array name is treated as a pointer.
- ▶ Example:

```
int find_largest(int a[], int n)
{
    int i, max;

    max = a[0];
    for (i = 1; i < n; i++)
        if (a[i] > max)
            max = a[i];
    return max;
}
```

- ▶ A call of `find_largest`:

```
largest = find_largest(b, N);
```

This call causes a pointer to the first element of **b** to be assigned to **a**; the array itself isn't copied.

Array Arguments (Revisited)

- ▶ The fact that an array argument is treated as a pointer has some important consequences.
- ▶ **Consequence 1:** When an ordinary variable is passed to a function, its value is copied; any changes to the corresponding parameter don't affect the variable.
- ▶ In contrast, **an array used as an argument isn't protected against change.**

Array Arguments (Revisited)

- ▶ For example, the following function modifies an array by storing zero into each of its elements:

```
void store_zeros(int a[], int n)
{
    int i;

    for (i = 0; i < n; i++)
        a[i] = 0;
}
```

Array Arguments (Revisited)

- ▶ To indicate that an array parameter won't be changed, we can include the word `const` in its declaration:

```
int find_largest(const int a[], int n)
{
    ...
}
```

- ▶ If `const` is present, the compiler will check that no assignment to an element of `a` appears in the body of `find_largest`.

Array Arguments (Revisited)

- ▶ **Consequence 2:** The time required to pass an array to a function **doesn't** depend on the size of the array.
- ▶ There's no penalty for passing a large array, since no copy of the array is made.

Array Arguments (Revisited)

- ▶ **Consequence 3:** An array parameter can be declared as a pointer if desired.
- ▶ `find_largest` could be defined as follows:

```
int find_largest(int *a, int n)
{
    ...
}
```
- ▶ Declaring `a` to be a pointer is equivalent to declaring it to be an array; the compiler treats the declarations as though they **were identical**.

Array Arguments (Revisited)

- ▶ Although declaring a *parameter* to be an array is the same as **declaring** it to be a pointer, **the same isn't true for a *variable*.**
- ▶ The following declaration causes the compiler to set aside space for 100 integers:

```
int a[100];
```

- ▶ The following declaration causes the compiler to allocate space for a pointer variable:

```
int *a;
```

Array Arguments (Revisited)

- ▶ In the latter case, a **is not** an array; attempting to use it as an array can have disastrous results.
- ▶ For example, the assignment
`*a = 0; /*** WRONG ***/`
will store 0 where a is pointing.
- ▶ Since we don't know where a is pointing, the effect on the program is undefined.

Array Arguments (Revisited)

- ▶ **Consequence 4:** A function with an array parameter can be passed an array “slice” – a sequence of consecutive elements.
- ▶ An example that applies `find_largest` to elements 5 through 14 of an array `b`:

```
largest = find_largest(&b[5], 10);
```

Using a Pointer as an Array Name

- ▶ C allows us to subscript a pointer as though it were an array name:

```
#define N 10  
...  
int a[N], i, sum = 0, *p = a;  
...  
for (i = 0; i < N; i++)  
    sum += p[i];
```

The compiler treats `p[i]` as `*(p+i)`.



In this lecture

- ▶ Types: data + operations
 - ▶ Pointers and arrays
 - ▶ Pointer arithmetic
 - ▶ Arrays as arguments
-
- ▶ Multi-dimensional arrays (pointers)

Extras:

- ▶ Variable-length arrays (pointers)




Pointers and Multidimensional Arrays

- ▶ Just as pointers can point to elements of one-dimensional arrays, they can also point to elements of multidimensional arrays.
- ▶ This section explores common techniques for using pointers to process the elements of multidimensional arrays.

Processing the Elements of a Multidimensional Array

- ▶ Recall that C stores two-dimensional arrays in row-major order.
- ▶ Layout of an array with r rows:



- ▶ If p initially points to the element in row 0, column 0, we can visit every element in the array by incrementing p repeatedly.

Processing the Elements of a Multidimensional Array

- ▶ Consider the problem of initializing all elements of the following array to zero:

```
int a[NUM_ROWS][NUM_COLS];
```

- ▶ The obvious technique would be to use nested `for` loops:

```
int row, col;  
...  
for (row = 0; row < NUM_ROWS; row++)  
    for (col = 0; col < NUM_COLS; col++)  
        a[row][col] = 0;
```

- ▶ If we view `a` as a one-dimensional array of integers, a single loop is sufficient:

```
int *p;  
...  
for (p = &a[0][0]; p <= &a[NUM_ROWS-1][NUM_COLS-1]; p++)  
    *p = 0;
```

Processing the Elements of a Multidimensional Array

- ▶ Although treating a two-dimensional array as one-dimensional may seem like cheating, it works with most C compilers.
- ▶ Techniques like this one definitely hurt program **readability**, but—at least with some older compilers—produce a compensating increase in **efficiency**.
- ▶ With many modern compilers, though, there's often little or no speed advantage.

Processing the Rows of a Multidimensional Array

- ▶ A pointer variable `p` can also be used for processing the elements in **just one row** of a two-dimensional array.
- ▶ To visit the elements of row `i`, we'd initialize `p` to point to element 0 in row `i` in the array `a`:

```
p = &a[i][0];
```

or we could simply write

```
p = a[i];
```

Processing the Rows of a Multidimensional Array

- ▶ For any two-dimensional array a , the expression $a[i]$ is a pointer to the first element in row i .
- ▶ To see why this works, recall that $a[i]$ is equivalent to $*(a + i)$.
- ▶ Thus,
 $\&a[i][0]$ is the same as $\&(*(a[i] + 0))$,
which is equivalent to
 $\&*a[i]$.
- ▶ This is the same as $a[i]$, since the $\&$ and $*$ operators cancel, in this case.

Processing the Rows of a Multidimensional Array

- ▶ A loop that clears row i of the array a :

```
int a[NUM_ROWS][NUM_COLS], *p, i;  
...  
for (p = a[i]; p < a[i] + NUM_COLS; p++)  
    *p = 0;
```

Processing the Rows of a Multidimensional Array

- ▶ Since `a[i]` is a pointer to row `i` of the array `a`, we can pass `a[i]` to a function that's expecting a one-dimensional array as its argument, i.e. a function expecting a one-dimensional array will also work with a row belonging to a two-dimensional array.
- ▶ Consider `find_largest`, it finds the largest element of a one-dimensional array. It can be used to determine the largest element in row `i` of the two-dimensional array `a`:

```
largest = find_largest(a[i], NUM_COLS);
```

Processing the Columns of a Multidimensional Array

- ▶ Processing the elements in a *column* of a two-dimensional array isn't as easy, because **arrays are stored by row**, not by column.
- ▶ A loop that clears **column *i*** of the array *a*:

```
int a[NUM_ROWS][NUM_COLS], (*p)[NUM_COLS], i;  
...  
for (p = &a[0]; p < &a[NUM_ROWS]; p++)  
    (*p)[i] = 0;
```

Using the Name of a Multidimensional Array as a Pointer

- ▶ The name of *any* array can be used as a pointer, regardless of how many dimensions it has, but some care is required.

- ▶ Example:

```
int a[NUM_ROWS][NUM_COLS];
```

a is *not* a pointer to `a[0][0]`; instead, it's a pointer to `a[0]`.

- ▶ C regards `a` as a one-dimensional array whose elements are one-dimensional arrays.
- ▶ When used as a pointer, `a` has type `int (*)[NUM_COLS]` (pointer to an integer array of length `NUM_COLS`).

Using the Name of a Multidimensional Array as a Pointer

- ▶ Knowing that **a points to a[0]** is useful for simplifying loops that process the elements of a two-dimensional array.
- ▶ Instead of writing

```
for (p = &a[0]; p < &a[NUM_ROWS]; p++)  
    (*p)[i] = 0;
```

to clear **column i** of the array a, we can write

```
for (p = a; p < a + NUM_ROWS; p++)  
    (*p)[i] = 0;
```

Using the Name of a Multidimensional Array as a Pointer

Careful about types, always! E.g.:

- ▶ We can “trick” a function into thinking that a multidimensional array is really one-dimensional.
- ▶ A first attempt at using using `find_largest` to find the largest element in `a`:

```
largest = find_largest(a, NUM_ROWS * NUM_COLS);  
/* WRONG */
```

This an error, because the type of `a` is **`int (*) [NUM_COLS]`** but `find_largest` is expecting an argument of type `int *`.

- ▶ The correct call:

```
largest = find_largest(a[0], NUM_ROWS * NUM_COLS);
```

`a[0]` points to element 0 in row 0, and it has type **`int *`** (after conversion by the compiler).



Pointers and Variable-Length Arrays (C99)

- ▶ Pointers are allowed to point to elements of variable-length arrays (VLAs).
- ▶ An ordinary pointer variable would be used to point to an element of a one-dimensional VLA:

```
void f(int n)
{
    int a[n], *p;
    p = a;
    ...
}
```

Pointers and Variable-Length Arrays (C99)

- ▶ When the VLA has more than one dimension, **the type of the pointer depends on the length of each dimension except for the first.**
- ▶ A two-dimensional example:

```
void f(int m, int n)
{
    int a[m][n], (*p)[n];
    p = a;
    ...
}
```

Since the type of `p` depends on `n`, which isn't constant, `p` is said to have a **variably modified type**.

Pointers and Variable-Length Arrays (C99)

- ▶ The validity of an assignment such as `p = a` can't always be determined by the compiler.
- ▶ The following code will compile but is correct only if `m` and `n` are equal:

```
int a[m][n], (*p)[m];  
p = a;
```

- ▶ If `m` is not equal to `n`, any subsequent use of `p` will cause undefined behavior.

Pointers and Variable-Length Arrays (C99)

- ▶ Variably modified types are subject to certain restrictions.
- ▶ The most important restriction: the declaration of a variably modified type must be **inside the body** of a function or in a function prototype.

Pointers and Variable-Length Arrays (C99)

- ▶ Pointer arithmetic works with VLAs.

- ▶ A two-dimensional VLA:

```
int a[m][n];
```

- ▶ A pointer capable of pointing to a row of a:

```
int (*p)[n];
```

- ▶ A loop that clears column *i* of a:

```
for (p = a; p < a + m; p++)  
    (*p)[i] = 0;
```