

Input/Output

Dr Deepayan Bhowmik Dr Andrea Bracciali

Introduction

- ▶ C's input/output library is the biggest and most important part of the **standard library**.
- ▶ The `<stdio.h>` header is the primary repository of input/output functions, including `printf`, `scanf`, `putchar`, `getchar`, `puts`, and `gets`.
- ▶ We've looked at some of these functions.
- ▶ ... and will introduce new functions, most of which deal with files.



Introduction

- ▶ In C99, some I/O functions belong to the `<wchar.h>` header. They deal with wide characters rather than ordinary characters.
- ▶ Functions in `<stdio.h>` that read or write data are known as **byte input/output functions**.
- ▶ Similar functions in `<wchar.h>` are called **wide-character input/output functions**.



Streams

- ▶ In C, the term **stream** means any **source of input** or any **destination for output**.
- ▶ Many programs obtain all their input from one stream (the **keyboard**) and write all their output to another stream (the **screen**).
- ▶ Larger programs may need additional streams.
- ▶ Streams **often represent files stored on various media**.



File Pointers

- ▶ Accessing a stream is done through a **file pointer**, which has type `FILE *`.
- ▶ The `FILE` type is declared in `<stdio.h>`.
- ▶ Certain streams are represented by file pointers with standard names.
- ▶ Additional file pointers can be declared as needed:

```
FILE *fp1, *fp2;
```



Standard Streams and Redirection

- ▶ `<stdio.h>` provides three standard streams:

<i>File Pointer</i>	<i>Stream</i>	<i>Default Meaning</i>
<code>stdin</code>	Standard input	Keyboard
<code>stdout</code>	Standard output	Screen
<code>stderr</code>	Standard error	Screen

- ▶ These streams are ready to use—we don't declare them, and we don't open or close them.



Text Files versus Binary Files

- ▶ `<stdio.h>` supports two kinds of files: text and binary.
- ▶ The bytes in a **text file** represent characters, allowing humans to examine or edit the file.
 - ▶ The **source** code for a C program is stored in a text file.
- ▶ In a **binary file**, bytes don't necessarily represent characters.
 - ▶ Groups of bytes might represent other types of data, such as integers and floating-point numbers.
 - ▶ An **executable** C program is stored in a binary file.



Text Files versus Binary Files

- ▶ Text files have two characteristics that binary files don't possess: carriage return (`'\r'`) and line feed (`'\n'`) – other encoding are used.
- ▶ **Text files are divided into lines.** Each line in a text file normally ends with one or two special characters.
 - ▶ Windows: carriage-return character (`'\r'`) followed by line-feed character (`'\n'`)
 - ▶ UNIX and newer versions of Mac OS: line-feed character
 - ▶ Older versions of Mac OS: carriage-return character

TRY IT!




Text Files versus Binary Files

- ▶ **Text files may contain a special “end-of-file” marker.**
 - ▶ In Windows, the marker is ' `\x1a` ' (Ctrl-Z), but it is not required.
 - ▶ Most other operating systems, including UNIX, have no special end-of-file character.
- ▶ **In a binary file, there are no end-of-line or end-of-file markers; all bytes are treated equally.**




Text Files versus Binary Files

- ▶ When data is written to a file, it can be stored in text form or in binary form.
- ▶ One way to store the number 32767 in a file would be to write it in text form as the **characters** 3, 2, 7, 6, and 7:



Text Files versus Binary Files

- ▶ The other option is to store the number in **binary**, which would take as few as two bytes:



- ▶ Storing numbers in binary can often save space.

Opening a File

- ▶ Opening a file for use as a stream requires a call of the `fopen()` function.
- ▶ **Prototype for `fopen()`:**

```
FILE *fopen(const char * filename,  
            const char * mode);
```

- ▶ `filename` is the name of the file to be opened.
 - ▶ This argument may include information about the file's location, such as a drive specifier or path.
- ▶ `mode` is a “mode string” that specifies what operations we intend to perform on the file.



Opening a File

- ▶ In Windows, **be careful** when the file name in a call of `fopen` includes **the `\` character**.

- ▶ The call

```
fopen("c:\project\test1.dat", "r")
```

will fail, because `\t` is treated as a character escape.

- ▶ One way to avoid the problem is to use `\\` instead of `\`:

```
fopen("c:\\project\\test1.dat", "r")
```

- ▶ An alternative is to use the `/` character instead of `\`:

```
fopen("c:/project/test1.dat", "r")
```



Opening a File

- ▶ `fopen` returns a file pointer that the program can (and usually will) save in a variable:

```
FILE* fp;  
fp = fopen("in.dat", "r");  
/* opens in.dat for reading */
```

- ▶ When it can't open a file, `fopen` returns a null pointer.



Modes

- ▶ Factors that determine which mode string to pass to `fopen`:
 - ▶ Which operations are to be performed on the file
 - ▶ Whether the file contains text or binary data



Modes

► Mode strings for text files:

<i>String</i>	<i>Meaning</i>
"r"	Open for reading
"w"	Open for writing (file need not exist)
"a"	Open for appending (file need not exist)
"r+"	Open for reading and writing, starting at beginning
"w+"	Open for reading and writing (truncate if file exists/create)
"a+"	Open for reading and writing (append if file exists)



Modes

► Mode strings for binary files:

<i>String</i>	<i>Meaning</i>
"rb"	Open for reading
"wb"	Open for writing (file need not exist)
"ab"	Open for appending (file need not exist)
"r+b" or "rb+"	Open for reading and writing, starting at beginning
"w+b" or "wb+"	Open for reading and writing (truncate if file exists)
"a+b" or "ab+"	Open for reading and writing (append if file exists)



Modes

- ▶ Note that there are different mode strings for **writing** data and **appending** data.
- ▶ When data is written to a file, it normally overwrites what was previously there.
- ▶ When a file is opened for appending, data written to the file is added at the end.



Closing a File

- ▶ The `fclose` function allows a program to close a file that it's no longer using.
- ▶ The argument to `fclose` must be a file pointer obtained from a call of `fopen`.
- ▶ `fclose` returns zero if the file was closed successfully.
- ▶ Otherwise, it returns the error code `EOF` (a macro defined in `<stdio.h>`).



Closing a File

- ▶ The outline of a program that opens a file for reading:

```
#include <stdio.h>
#include <stdlib.h>

#define FILE_NAME "example.dat"

int main(void)
{
    FILE *fp;

    fp = fopen(FILE_NAME, "r");
    if (fp == NULL) {
        printf("Can't open %s\n", FILE_NAME);
        exit(EXIT_FAILURE);
    }
    ...
    fclose(fp);
    return 0;
}
```



A Note on opening a File

- ▶ It's not unusual to see the call of `fopen` combined with the declaration of `fp`:

```
FILE *fp = fopen(FILE_NAME, "r");
```

or the test against `NULL`:

```
if ((fp = fopen(FILE_NAME, "r")) == NULL) ...
```



Obtaining File Names from the Command Line

- ▶ Ways to supply file names to a program:
 - ▶ Building file names into the program **doesn't provide much flexibility**.
 - ▶ Prompting the user to enter file names can be awkward.
 - ▶ Having the program obtain file names from the **command line** is often the best solution.
- ▶ An example that uses the command line to supply two file names to a program named `demo`:

```
demo names.dat dates.dat
```
- ▶ In Netbeans, Project Properties > Run > Run Command, eg.

```
"${OUTPUT_PATH}" arg1 arg2
```



Obtaining File Names from the Command Line

- ▶ Earlier lecture showed how to access command-line arguments by defining `main` as a function with two parameters:

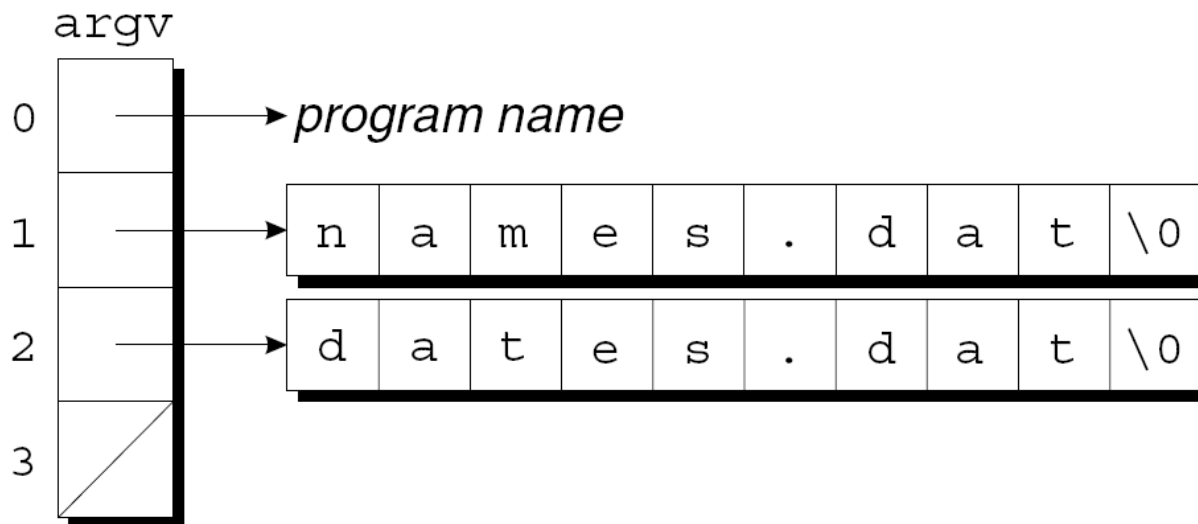
```
int main(int argc, char *argv[])  
{  
    ...  
}
```

- ▶ `argc` is the number of command-line arguments.
- ▶ `argv` is an array of pointers to the argument strings.



Obtaining File Names from the Command Line

- ▶ `argv[0]` points to the program name, `argv[1]` through `argv[argc-1]` point to the remaining arguments, and `argv[argc]` is a null pointer.
- ▶ In the demo example, `argc` is 3 and `argv` has the following appearance:



Program: Checking Whether a File Can Be Opened

- ▶ The `canopen.c` program determines if a file exists and can be opened for reading.
- ▶ The user will give the program a file name to check:
`canopen file`
- ▶ The program will then print either `file` can be opened or `file` can't be opened.
- ▶ If the user enters the wrong number of arguments on the command line, the program will print the message `usage: canopen filename.`



canopen.c

```
/* Checks whether a file can be opened for reading */
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    FILE *fp;

    if (argc != 2) {
        printf("usage: canopen filename\n");
        exit(EXIT_FAILURE);
    }

    if ((fp = fopen(argv[1], "r")) == NULL) {
        printf("%s can't be opened\n", argv[1]);
        exit(EXIT_FAILURE);
    }

    printf("%s can be opened\n", argv[1]);
    fclose(fp);
    return 0;
}
```



Formatted I/O

- ▶ The next group of library functions use format strings to control reading and writing.
- ▶ `printf` and related functions are able to convert data from numeric form to character form during output.
- ▶ `scanf` and related functions are able to convert data from character form to numeric form during input.



The `..printf` Functions

- ▶ The `fprintf` and `printf` functions write a variable number of data items to an output stream, using a format string to control the appearance of the output.
- ▶ The prototypes for both functions end with the `...` symbol (an ***ellipsis***), which indicates a variable number of additional arguments:

```
int fprintf(FILE * restrict stream,  
            const char * restrict format, ...);  
int printf(const char * restrict format, ...);
```

- ▶ Both functions return the number of characters written; a **negative return** value indicates that an error occurred.



The `..printf` Functions

- ▶ `printf` **always writes to stdout, whereas `fprintf` writes to the stream indicated by its first argument:**

```
printf("Total: %d\n", total);  
    /* writes to stdout */  
fprintf(fp, "Total: %d\n", total);  
    /* writes to fp */
```

- ▶ **A call of `printf` is equivalent to a call of `fprintf` with `stdout` as the first argument**



TRY IT



The `...scanf` Functions

- ▶ `fscanf` and `scanf` read data items from an input stream, using a format string to indicate the layout of the input.
- ▶ After the format string, any number of pointers—each pointing to an object—follow as additional arguments.
- ▶ Input items are **converted** (according to conversion specifications in the format string) and stored in these objects.



The **...scanf** Functions

- ▶ **scanf** always reads from `stdin`, whereas **fscanf** reads from the stream indicated by its first argument:

```
scanf("%d%d", &i, &j);  
    /* reads from stdin */  
fscanf(fp, "%d%d", &i, &j);  
    /* reads from fp */
```

- ▶ A call of **scanf** is equivalent to a call of **fscanf** with `stdin` as the first argument.



Character I/O

- ▶ Library functions to read and write single characters.
- ▶ These functions work equally well with text streams and binary streams.
- ▶ The functions treat characters as values of `type int`, not `char`.



Character I/O

- ▶ Library functions to read and write single characters.
- ▶ `putchar` **writes one character to the stdout stream:**
`putchar(ch); /* writes ch to stdout */`
- ▶ `fputc` **write a character to an arbitrary stream:**
`fputc(ch, fp); /* writes ch to fp */`



Input Functions

- ▶ `getchar` reads a character from `stdin`:

```
ch = getchar();
```

- ▶ `fgetc` read a character from an arbitrary stream:

```
ch = fgetc(fp);
```

The function treats the character as **an unsigned char value** (which is then converted to **int type** before it's returned).

- ▶ As a result, they never return a negative value other than EOF.



Program: Copying a File

- ▶ The `fcopy.c` program makes a copy of a file (code next!).
- ▶ The names of the original file and the new file will be specified on the command line when the program is executed.
- ▶ An example that uses `fcopy` to copy the file `f1.c` to `f2.c`:

```
fcopy f1.c f2.c
```

- ▶ `fcopy` will **issue an error message** if there aren't exactly two file names on the command line or if either file can't be opened.



Program: Copying a File

- ▶ Using `"rb"` and `"wb"` as the file modes enables `fcopy` to copy **both text and binary** files.
- ▶ If we used `"r"` and `"w"` instead, the program wouldn't necessarily be able to copy binary files.



fcopy.c

```
/* Copies a file */

#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    FILE *source_fp, *dest_fp;
    int ch;

    if (argc != 3) {
        fprintf(stderr, "usage: fcopy source dest\n");
        exit(EXIT_FAILURE);
    }
}
```



```
if ((source_fp = fopen(argv[1], "rb")) == NULL) {  
    fprintf(stderr, "Can't open %s\n", argv[1]);  
    exit(EXIT_FAILURE);  
}  
  
if ((dest_fp = fopen(argv[2], "wb")) == NULL) {  
    fprintf(stderr, "Can't open %s\n", argv[2]);  
    fclose(source_fp);  
    exit(EXIT_FAILURE);  
}  
  
while ((ch = fgetc(source_fp)) != EOF)  
    fputc(ch, dest_fp);  
  
fclose(source_fp);  
fclose(dest_fp);  
return 0;  
}
```



Line I/O

- ▶ Library functions in the next group are able to read and write lines.
- ▶ These functions are **used mostly with text streams**, although it's legal to use them with binary streams as well.



Output Functions

- ▶ The `puts` function writes a string of characters to `stdout`:

```
puts("Hi, there!"); /* writes to stdout */
```

- ▶ After it writes the characters in the string, `puts` always adds a **new-line character**.



Output Functions

- ▶ `fputs` is a more general version of `puts`.

- ▶ Its second argument indicates the stream to which the output should be written:

```
fputs("Hi, there!", fp); /* writes to fp */
```

- ▶ Unlike `puts`, the `fputs` function doesn't write a new-line character unless one is present in the string.
- ▶ Both functions **return EOF** if a write error occurs; otherwise, they return a nonnegative number.



Input Functions

- ▶ The `gets` function reads **a line of input** from `stdin`:
`gets(str); /* reads a line from stdin */`
- ▶ `gets` reads characters one by one, **storing them in the array** pointed to by `str`, until it reads a new-line character (which it discards).
- ▶ `fgets` is a more general version of `gets` that can read from any stream.
- ▶ `fgets` **is also safer than** `gets`, since it limits the number of characters that it will store.



Input Functions

- ▶ A call of `fgets` that reads a line into a character array named `str`:

```
fgets(str, sizeof(str), fp);
```
- ▶ `fgets` will read characters until it reaches the first new-line character or `sizeof(str) - 1` characters have been read.
- ▶ If it reads the new-line character, `fgets` stores it along with the other characters.



Input Functions

- ▶ Both `gets` and `fgets` **return a null pointer** if a read error occurs or they reach the end of the input stream before storing any characters.
- ▶ Otherwise, both return their first argument, which **points to the array** in which the input was stored.
- ▶ Both functions store a **null character** at the end of the string.



Block I/O

- ▶ The `fread` and `fwrite` functions allow a program to read and write large blocks of data in a single step.
- ▶ `fread` and `fwrite` are used primarily with binary streams, although—with care—it's possible to use them with text streams as well.



Block I/O

- ▶ `fwrite` is designed to copy **an array from memory to a stream**.
- ▶ Arguments in a call of `fwrite`:
 - ▶ Address of array
 - ▶ Size of each array element (**in bytes**)
 - ▶ Number of elements to write
 - ▶ File pointer
- ▶ A call of `fwrite` that writes the entire contents of the array `a`:

```
fwrite(a, sizeof(a[0]), sizeof(a) / sizeof(a[0]), fp);
```



Block I/O

- ▶ `fwrite` returns the number of elements actually written.
- ▶ This number will be less than the third argument if a write error occurs.



Block I/O

- ▶ `fread` will read the elements of an **array from a stream**.
- ▶ A call of `fread` that reads the contents of a file into the array `a`:

```
n = fread(a, sizeof(a[0]),  
          sizeof(a) / sizeof(a[0]), fp);
```

- ▶ `fread`'s **return value** indicates **the actual number of elements read**.
- ▶ This number should equal the third argument unless the end of the input file was reached or a read error occurred.



Block I/O

- ▶ `fwrite` is convenient for a program that needs to store data in a file before terminating.
- ▶ Later, the program (or another program) can use `fread` to read the data back into memory.
- ▶ The data doesn't need to be in array form.
- ▶ A call of `fwrite` that writes a structure variable `s` to a file:

```
fwrite(&s, sizeof(s), 1, fp);
```



String I/O

- ▶ The next functions can read and write data **using a string as though it were a stream.**
- ▶ `sprintf` and `snprintf` write characters into a string.
- ▶ `sscanf` reads characters from a string.



Output Functions

- ▶ The `sprintf` function writes output into a character array (pointed to by its first argument) instead of a stream.
- ▶ A call that writes "9/20/2010" into `date`:

```
sprintf(date, "%d/%d/%d", 9, 20, 2010);
```
- ▶ `sprintf` adds a null character at the end of the string.
- ▶ It returns the number of characters stored (not counting the null character).



Output Functions

`sprintf`

- ▶ can be used to format data, with the result saved in a string until it's time to produce output.
- ▶ it is also convenient for converting numbers to character form.



Output Functions

- ▶ The `snprintf` function (new in C99) is the same as `sprintf`, except for an additional second parameter named `n`.
- ▶ **No more than `n - 1` characters will be written** to the string, not counting the terminating null character, which is always written unless `n` is zero.
- ▶ **Example:**

```
snprintf(name, 13, "%s, %s", "Einstein", "Albert");
```

The string "Einstein, Al" is written into `name`.



Output Functions

- ▶ `snprintf` returns the number of characters that would have been written (not including the null character) had there been no length restriction.
- ▶ If an encoding error occurs, `snprintf` returns a negative number.
- ▶ To see if `snprintf` had room to write all the requested characters, we can test whether its return value was nonnegative and less than `n`.



Input Functions

- ▶ The `sscanf` function is similar to `scanf` and `fscanf`.
- ▶ `sscanf` reads **from a string** (pointed to by its first argument) instead of reading from a stream.
- ▶ `sscanf`'s second argument is **a format string** identical to that used by `scanf` and `fscanf`.



Input Functions

- ▶ `sscanf` is handy for **extracting data from a string** that was read by another input function.
- ▶ An example that uses `fgets` to obtain a line of input, then passes the line to `sscanf` for further processing:

```
fgets(str, sizeof(str), stdin);  
    /* reads a line of input */  
sscanf(str, "%d%d", &i, &j);  
    /* extracts two integers */
```



